




MC training progress on cart-pole task


100-episode return (running mean)



Episode length (running mean)



Behavior entropy (running mean)



Sample loss (RMSE) (running mean)

